

See discussions, stats, and author profiles for this publication at: <https://www.researchgate.net/publication/323149800>

# Speech Recognition for the iCub Platform

Article · February 2018

DOI: 10.3389/frobt.2018.00010

CITATIONS

0

READS

7

4 authors, including:



**Bertrand Higy**

Istituto Italiano di Tecnologia

2 PUBLICATIONS 5 CITATIONS

[SEE PROFILE](#)



**Alessio Mereta**

European Space Agency

6 PUBLICATIONS 14 CITATIONS

[SEE PROFILE](#)



**Leonardo Badino**

Istituto Italiano di Tecnologia

45 PUBLICATIONS 322 CITATIONS

[SEE PROFILE](#)

Some of the authors of this publication are also working on these related projects:



ALLSpeak [View project](#)



Differential Intelligence [View project](#)

All content following this page was uploaded by **Bertrand Higy** on 26 February 2018.

The user has requested enhancement of the downloaded file.

# Speech recognition for the iCub platform

Bertrand Higy<sup>1,2</sup>, Alessio Mereta<sup>3</sup>, Giorgio Metta<sup>1</sup> and Leonardo Badino<sup>4,\*</sup>

<sup>1</sup>*iCub Facility, Istituto Italiano di Tecnologia, Genova, Italy*

<sup>2</sup>*Università di Genova, Genova, Italy*

<sup>3</sup>*Advanced Concepts Team, European Space Agency, Noordwijk, The Netherlands*

<sup>4</sup>*Center for Translational Neurophysiology of Speech and Communication, Istituto Italiano di Tecnologia, Ferrara, Italy*

Correspondence\*:

Leonardo Badino

leonardo.badino@iit.it

2 Figures: 2

3 Tables: 1

4 Word count: 2984

## 5 ABSTRACT

6 This paper describes open source software (available at <https://github.com/robotology/natural-speech>) to build automatic speech recognition (ASR) systems and run them within the YARP  
7 platform. The toolkit is designed (i) to allow non-ASR experts to easily create their own ASR  
8 system and run it on iCub, and (ii) to build deep learning-based models specifically addressing  
9 the main challenges an ASR system faces in the context of verbal human-iCub interactions. The  
10 toolkit mostly consists of Python, C++ code and shell scripts integrated in YARP.

11 As additional contribution, a second codebase (written in Matlab) is provided for more expert  
12 ASR users who want to experiment with bio-inspired and developmental learning-inspired ASR  
13 systems. Specifically, we provide code for two distinct kinds of speech recognition: “articulatory”  
14 and “unsupervised” speech recognition. The first is largely inspired by influential neurobiological  
15 theories of speech perception which assume speech perception to be mediated by brain motor  
16 cortex activities. Our articulatory systems have been shown to outperform strong deep learning-  
17 based baselines. The second type of recognition systems, the “unsupervised” systems, do  
18 not use any supervised information (contrary to most ASR systems, including our articulatory  
19 systems). To some extent, they mimic an infant who has to discover the basic speech units of a  
20 language by herself.

21 In addition, we provide resources consisting of pre-trained deep learning models for ASR, and a  
22 2,5-hours speech dataset of spoken commands, the VoCub dataset, which can be used to adapt  
23 an ASR system to the typical acoustic environments in which iCub operates.

24  
25 **Keywords:** automatic speech recognition, yarp, tensorflow, code:python, code:matlab, code:C++

## 1 INTRODUCTION

26 Several applications use speech to give instructions to iCub, often relying on proprietary software. However,  
27 the robot operates in specific conditions where those systems may perform poorly. An open and easy-to-use  
28 system that would reliably recognize commands in this context would thus be a very desirable tool. We

present here a first codebase, henceforth *iCubRec*, which has been built to provide such services to the community of iCub users. It allows to train and run state-of-the-art deep neural network (DNN)-based automatic speech recognition (ASR).

As an additional contribution, a second codebase, henceforth *bioRec*, allows to experiment with novel DNN-based recognition systems that share the same bio-inspired and developmental learning view that gave birth to iCub (Lungarella et al., 2003). *bioRec* is self-contained and independent of *iCubRec*, however its DNN-based acoustic models can effortlessly be used within *iCubRec*.

Finally, in addition to the code, we are also providing resources to facilitate the implementation of a command recognizer: (i) the VoCub dataset, a dataset of registered vocal commands, (ii) pretrained Gaussian Mixture Model (GMM)- and DNN-based acoustic models to perform recognition.

Our code, as well as the resources, are released under GPLv3 license. The code is available at <https://github.com/robotology/natural-speech> (doi:10.5281/zenodo.1064043).

## 2 ICUBREC

### 2.1 Application and utility

An ASR system for iCub typically operates in challenging conditions. We have identified three specific factors which we want the system to be robust to:

- noise; the robot often operates in noisy environments (e.g., noisy servers and computers running, concurrent speakers, the robot itself generating noise).
- accents; the teams working with iCub are international and the robot needs to recognize spoken commands uttered with a wide variety of foreign accents.
- distance and movement; distant speech recognition is an important research topic in ASR and has been the focus of many recent challenges (e.g., the Chime4 challenge<sup>1</sup>). When the speaker-microphone distance increases, the speech signal-to-noise ratio decreases and signal distortions due to reverberation (in indoor environments) increases. A non-fixed distance, due to a moving speaker and/or microphone, adds further complexity to the task.

Although deep learning has recently produced excellent results in ASR, it still suffers the training-testing mismatched conditions problem. Proprietary ASR systems may perform poorly in the aforementioned acoustic/speech conditions mainly because such conditions are not well covered by their training datasets. We have addressed this problem by building a dataset (VoCub dataset) that covers such conditions and by providing tools to easily adapt a DNN to it.

Other than robust, an ASR system for iCub should be easy-to-use, open and modular. Usability is necessary to allow all iCub mindware developers, who mostly have no ASR background, to train and run ASR on iCub. For this reason we provide pretrained GMM- and DNN-based acoustic models that can be used out of the box with the existing code. At the same time, we want more advanced users to easily modify and adapt the code to their own needs. This can only be done if everything is open and well modularized.

---

<sup>1</sup> [http://spandh.dcs.shef.ac.uk/chime\\_challenge/chime2016/](http://spandh.dcs.shef.ac.uk/chime_challenge/chime2016/)

## 63 2.2 Methods

64 To facilitate the understanding of the *iCubRec* module for non-ASR experts we provide here the definition  
 65 of few basic ASR terms. A standard ASR system consists of 4 main parts: an acoustic feature extraction  
 66 step which extracts spectral features from the input acoustic waveform; an acoustic model which relates  
 67 the extracted features to sub-words (e.g, phonemes, such as consonants and vowels) and then words (i.e.,  
 68 computes the likelihood that vectors of features are generated by a candidate word); a language model,  
 69 which is independent of the acoustic signals and incorporates prior knowledge about a specific language  
 70 (e.g., the probability that the word “barks” follows the word “dog”); and a speech decoder which performs  
 71 word recognition by computing the most probable sequence of words of the utterance, given: a) the acoustic  
 72 model; b) the language model; c) the dictionary, which consists of all words the system has to recognize  
 73 along with their phoneme transcriptions. Acoustic modeling is usually done using a Hidden Markov Model  
 74 (HMM) which is well suited for sequential data like speech. HMMs combine transition probabilities (i.e.,  
 75  $p(s_t|s_{t-1})$  where  $s_t$  is a phone label at time  $t$ ) with observation probabilities (i.e.,  $p(o_t|s_t)$ , where  $o_t$  is  
 76 the input vector of acoustic feature at time  $t$ ). The core difference between classical GMM-HMM vs.  
 77 hybrid DNN-HMM acoustic models simply resides on whether GMMs or DNNs are used to compute the  
 78 observation probabilities.

## 79 2.3 Code description

80 iCubRec code is based on the Hidden Markov Model Toolkit (HTK) (Young et al., 2015). However, as  
 81 the training capabilities for DNNs are still quite limited in HTK, we also consider the alternative possibility  
 82 to train a network with Tensorflow (Martín Abadi et al., 2015) and convert it to HTK format for use in  
 83 decoding. Although in the later case the DNN is still restricted to the architectures recognized by HTK (for  
 84 now, only feedforward networks with a limited set of activation functions), this gives more flexibility and  
 85 control over the training process. Additionally, the use of Tensorflow allows to easily adapt a pretrained  
 86 DNN to new adaptation data

87 The code consists of scripts for:

- 88 • acoustic model training with GMMs
- 89 • acoustic model training with DNNs
- 90 • speech decoding
- 91 • integration within YARP for online speech decoding.

92 iCubRec is a combination of Python 3, Perl and shell scripts, and was written for HTK 3.5 and Tensorflow  
 93 1.0.

### 94 2.3.1 GMM-based acoustic modeling

95 Before the advent of DNNs, GMM-HMM systems were state of the art for acoustic modeling in speech  
 96 recognition. Although they are significantly outperformed by neural networks (Seltzer et al., 2013; Dahl  
 97 et al., 2012), GMMs are still widely used if only to compute the phone labels / speech segments alignments  
 98 needed to train a DNN (Dahl et al., 2012). The folder `gmm_training` provides a set of scripts to train  
 99 GMM-HMMs using HTK. These scripts are based on Keith Vertanen’s code (Vertanen, 2006) and allow to  
 100 build models similar to the ones described by Woodland et al. (1994). The recipe is originally intended for  
 101 TIMIT (Garofolo et al., 1993b) and Wall Street Journal (WSJ) (Garofolo et al., 1993a) datasets and has  
 102 been adapted for the Chime4 challenge (Vincent et al., 2016) and VoCub datasets.

### 2.3.2 DNN-based acoustic modeling

Once the speech signal has been aligned (presumably using GMM-HMMs), a DNN-based model can be trained. Two alternatives are available: (i) using the scripts in `dnn_training/htk` to train a model with HTK or (ii) using the code under `dnn_training/tf` to train the net with Tensorflow. The scripts proposed here are currently restricted to TIMIT and WSJ, but support for additional datasets will be added soon.

### 2.3.3 Speech decoding

With a model trained with HTK (GMM-based or DNN-based), it is then straightforward to perform recognition on a new utterance. The folder `offline_decoding` provides an example of decoding on pre-recorded data with HTK. Additionally, `export_for_htk.py` shows how to easily extract the parameters of a net trained with Tensorflow and convert them into HTK format.

### 2.3.4 Integration with YARP

All the code presented so far is meant to train and test a system offline. `yarp_decoding` folder provides the modules necessary to use an existing model within YARP and perform online recognition. A streaming service based on `yarp.js`<sup>2</sup> allows to record sound from any device equipped with a microphone and a web browser. Two other modules are provided: `rcrld_yarphear_asr` which saves the recorded data in a file, and the `decoder` (based on HVite tool from HTK) for feature extraction and command decoding. The application `speechrec.xml` is available to easily run and connect all the modules.

## 2.4 Resources

### 2.4.1 The VoCub dataset

Recording a dataset has two main advantages: (i) it allows to easily test the recognition system and reliably estimate its performance in real conditions, and (ii) can be used to adapt the system in order to reduce the training/testing mismatch problem. For this reasons we have recorded examples of the commands we want to recognize within real-usage scenarios. That resulted in the VoCub dataset<sup>3</sup>.

The recordings consist of spoken English commands addressed to iCub. There are 103 unique commands (see Table 1 for some examples), composed of 62 different words. We recorded 29 speakers, 16 males and 13 females, 28 of them are non-native English speakers. We finally obtained 118 recordings from each speaker: of the 103 unique commands, 88 were recorded once, and 15 twice (corresponding to sentences containing rare words). This results in about 2 hours and 30 minutes of recording in total.

A split of the speakers into training, validation and test sets is proposed with 21, 4 and 4 speakers per set respectively. The files are organized with the following convention `setid/spkrid/spkrid_cond_recid.wav`, where:

- `setid` identifies the set: `tr` for training, `dt` for validation and `et` for testing.
- `spkrid` identifies the speaker: from 001 to 021 for training, 101 to 104 for validation and 201 to 204 for testing.
- `cond` identifies the condition (see below).
- `recid` identifies the record within the condition (starting from 0 and increasing).

<sup>2</sup> <https://github.com/robotology/yarp.js>

<sup>3</sup> freely available at <https://robotology.github.io/natural-speech/vocub/>

**Table 1.** 10 examples of the commands used in the VoCub dataset

I will teach you a new object.
This is an octopus.
What is this?
Let me show you how to reach the car with your left arm.
Let me show you how to reach the turtle with your right arm.
There you go.
Grasp the ladybug.
Where is the car?
No, here it is.
See you soon.

140 The commands were recorded in two different conditions, a non-static (`cond = 1`) and a static condition  
 141 (`cond = 2`), with an equal number of recorded utterances per condition.

142 In the static condition, the speaker sat in front of two screens where the sentences to read were displayed.  
 143 In the non-static condition, the commands were provided to the subject verbally through a speech synthesis  
 144 system, and the subject had to repeat them while performing a secondary manual task. This secondary task  
 145 was designed to be simple enough to not impede the utterance repetition task, while requiring people to  
 146 move around the robot. The distance between the speaker and the microphone in this last condition ranges  
 147 from 50 cm to 3 m.

148 We also registered a set of additional sentences for the testing group (same structure but different  
 149 vocabulary) to test the recognition system for new commands not seen during training. The sentences  
 150 consist of 20 new commands, pronounced by each speaker of the test set twice: once in non-static condition  
 151 (`cond = 3`) and once in static condition (`cond = 4`).

## 152 2.4.2 Trained models

153 As not all the datasets used in our scripts are freely available, and in order to ease the use of our system,  
 154 we provide pre-trained acoustic models that can be used out of the box. The `models/README.md`  
 155 file contains links to download GMM-based models trained on WSJ, Chime4 and VoCub datasets, and  
 156 DNN-based models trained on TIMIT and WSJ. Additional DNN-based models will be added in the future.  
 157 Further details about the different models and the precise training procedure can be found in the same file.

## 158 2.5 Example of use

159 A good demonstration of the capabilities of the code presented so far is given in the file  
 160 [icubrec/DEMO.md](#). In a few simple steps, the user is shown how to perform offline decoding on  
 161 the VoCub dataset with a pretrained model. This example is accessible to novice ASR users and does not  
 162 require any proprietary dataset.

163 A more in-depth example is given in `icubrec/TUTORIAL.md`, which provides detailed instruction on  
 164 how to train a full ASR system on the WSJ dataset. This tutorial goes through all the main steps: training  
 165 of a GMM-based acoustic model, computation of the alignments, training of a DNN-based acoustic model  
 166 using those alignments, and finally decoding of the test sentences.



### 3 BIOREC

#### 3.1 Application and utility

Our module for bio- and cognitive science-inspired ASR is composed of two distinct parts serving different purposes: *Articulatory Phone Recognition* and *Unsupervised/Developmental ASR*.

##### 3.1.1 Articulatory Phone Recognition

This part includes modules `phonerec` and `pce_phonerec`, which build *articulatory* phone recognition systems. A phone recognition system recognizes the sequence of phones of an utterance. It can roughly be identified as an ASR system without language model and dictionary. *Articulatory* phone recognition uses prior information about how the vocal tract moves when producing speech sounds. This *articulatory* view is strongly motivated by influential neurobiological theories of speech perception that assume a contribution of the brain motor cortex to speech perception (Pulvermüller and Fadiga, 2010), and have been shown to outperform strong DNN-based baselines where no prior articulatory information is used (see e.g., Badino et al. (2016)).

##### 3.1.2 Unsupervised/Developmental ASR

The second part of *bioRec*, `zerorchallenge`, builds “unsupervised” ASR systems. Most recognition systems, including the *articulatory* systems, are trained on supervised data, where training utterances are associated to phonetic transcriptions, and the inventory of phones is given. This learning setting is far easier than the learning setting of an infant who has to acquire her native language and has to discover the basic units of the language on her own. In order to better understand how an infant can acquire the phone inventory during development from raw “unsupervised” utterances, we have created “unsupervised” ASR systems that were submitted and evaluated at the 1st Zero Resource Speech Challenge (ZRS challenge) (Versteegh et al., 2015).

#### 3.2 Methods

##### 3.2.1 Articulatory Phone Recognition

The *articulatory* phone recognition module consists of 2 parts depending on how speech production information is represented:

- `phonerec`; speech production is represented in the form of actual measurements of vocal tract movements, collected through instruments such as the electromagnetic articulograph (Richmond et al., 2011);
- `pce_phonerec`; vocal tract movements are initially described by discrete linguistic features and actual measurements are not used.

`phonerec`. In this module, prior information of speech production is built by learning, during training, an acoustic-to-articulatory mapping that allows to recover vocal tract movements, i.e., reconstructed articulatory features (AFs), from the acoustic signal (Badino et al., 2012, 2016). The reconstructed AFs are then appended to the usual input acoustic vector of the DNN that computes phone state posterior probabilities, i.e, the acoustic model DNN (see Figure 1, which shows the simplest strategy). Additionally, our code allows to apply autoencoder (AE)-based transformations to the original AFs in order to improve performance. AEs are a special kind of DNN that attempts to reconstruct its input after encoding it, typically through a lossy encoding. More details and evaluation results can be found in (Badino et al., 2016).

205 `pce_phonerec`. In this module, AFs are derived (through a DNN) from linguistic discrete features  
 206 (referred to as phonetic context embedding). They are used as secondary target for the acoustic model  
 207 DNN within a multi-task learning (MTL) strategy (Caruana, 1997). This strategy forces the DNN to learn  
 208 a motor representation without the need for time-consuming collection of actual articulatory data. Our  
 209 approach outperforms strong alternative MTL-based approaches (Badino, 2016).

### 210 3.2.2 Unsupervised/Developmental ASR

211 `zerorchallenge` is the module building the unsupervised/developmental ASR systems we submitted  
 212 to Task1 of the ZRS challenge at Interspeech 2015 (Versteegh et al., 2015). The goal of the challenge was  
 213 to compare systems that create new acoustic representations that can discriminate examples of minimal  
 214 pairs, i.e., words differing only in one phoneme (e.g., “hat” vs. “had”), while identifying as a single entity  
 215 different examples of a same word. Specifically, we focused on extracting discrete/symbolic representations,  
 216 which equals to automatically discovering the inventory of (phone-like) sub-words of a language. Our  
 217 core strategy is based on AEs (Badino et al., 2014), as shown in Figure 2. The provided scripts build 2  
 218 novel systems, one based on binarized AEs and one on Hidden Markov Model Encoders (HMM-Encoders)  
 219 (Badino et al., 2015).

220 A binarized AE is an AE whose encoding layer nodes are binary. At each time step, it transforms a vector  
 221 of real-valued acoustic features into a vector of binary units which in turn is associated to a positive integer  
 222 corresponding to a discovered specific sub-word.

223 The HMM-Encoder combines an AE with a HMM<sup>4</sup>. An approach solely based on AEs ignores the  
 224 sequential nature of speech and inter-subword dependencies. The HMM-Encoder was proposed to  
 225 specifically address these potential weaknesses.

## 226 3.3 Code description and example of use

227 All code is written in Matlab and uses the Parallel Processing Toolbox to allow fast DNN training with  
 228 GPUs. All modules were tested in Matlab 2013a and 2015a.

### 229 3.3.1 Articulatory Phone Recognition


230 `phonerec`. The file `ploclassify.m` allows to train and test articulatory phone recognition systems.  
 231 It requires the `inivar.m` configuration file where it is possible to define, e.g.: the type of AFs through  
 232 `cmotortype` (e.g., AE-transformed AFs or “plain” AFs), the hyperparameters of the acoustic model  
 233 DNN (`parnet_classifier`), and of the acoustic-to-articulatory mapping DNN (`parnet_regress`).

234 The folder `demo` contains 2 examples to build and evaluate a baseline (`audio1_motor0_rec0`) and an  
 235 articulatory phone recognition system (`audio1_motor3_rec1`) on the `mngu0` dataset (Richmond et al.,  
 236 2011). The dataset used here (available at [https://zenodo.org/record/836692/files/bioRec\\_Resources.tar.gz](https://zenodo.org/record/836692/files/bioRec_Resources.tar.gz),  
 237 under `/bioRec_Resources/phonerec_mngu0/`) is a preprocessed version of the `mngu0` dataset.

238 `pce_phonerec`. This articulatory phone recognition system is trained and evaluated by running  
 239 `mtkpr_pce.m`. It can be compared with an alternative MTL based strategy proposed by Microsoft  
 240 researchers (Seltzer and Droppo, 2013), by running the script `mtkpr_baseline.m`. All systems are  
 241 trained and tested on the TIMIT dataset, which unfortunately is not freely available. Training on different  
 242 datasets would require some small dataset-dependent modifications to the look-up table used to extract  
 243 discrete linguistic features from phone names.

<sup>4</sup> Our HMM training code is a modified version of code from K. Murphy’s BayesianNet toolbox, available at <https://github.com/bayesnet/bnt>





**Figure 1.** An example of articulatory phone recognition. Here the simplest strategy available in phonerec is shown.  $o_t$  is a vector of acoustic features, while  $s_t$  is a phone state.

244 We have created a Python+Tensorflow implementation the DNN training proposed in this module which  
 245 will be soon available.

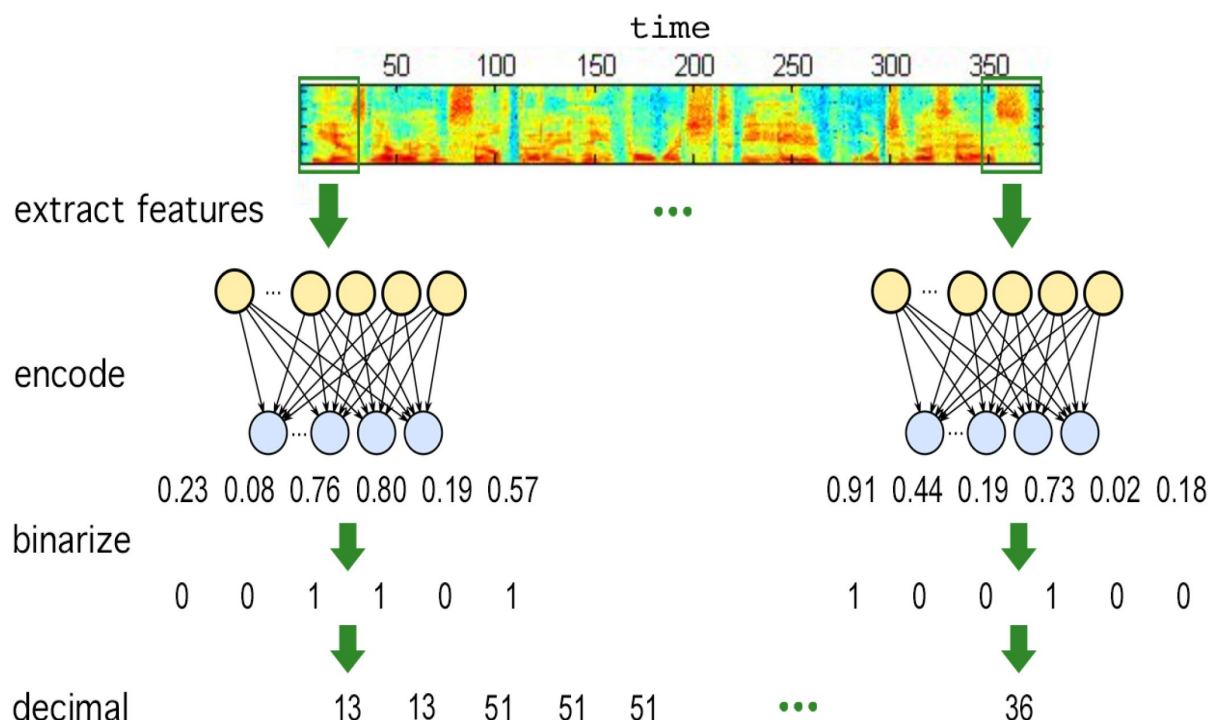
### 246 3.3.2 Unsupervised/Developmental ASR

247 We provide scripts that receive as input one of the datasets provided by the ZRS challenge,  
 248 train one of the unsupervised ASR systems (on the training utterances), and return the testing  
 249 utterances in a new discrete representation with a positive integer at each time step. We  
 250 additionally provide the 3 datasets from the ZRS challenge already transformed to be processed  
 251 by our scripts (available at [https://zenodo.org/record/836692/files/bioRec\\_Resources.tar.gz](https://zenodo.org/record/836692/files/bioRec_Resources.tar.gz), under  
 252 /bioRec\_Resources/zerochallenge/). The output format allows to evaluate the output file with  
 253 the tools provided for the challenge (Versteegh et al., 2015).

### 254 3.3.3 Utilities

255 All utilities used by the phonerec, pce\_phonerec and zerorchallenge are in:

- 256 • `netutils`. Contains functions to train and run DNNs, e.g.: standard DNN training, Deep Belief  
 257 Network-based DNN pretraining (Hinton et al., 2006), MTL training, DNN forward pass (i.e., to  
 258 evaluate a DNN), deep auto-encoder training, including training of some AEs we have recently  
 259 proposed specifically for speech.



**Figure 2.** Overview of the AE-based approach to subword learning.

- `utils`. This folder contains all utilities that do not pertain to DNNs. These include: data loading and normalization, phone language models computation, Viterbi-based phone decoding, phone error rate computation and analysis of error.

## 4 CONCLUSION

In this paper we have described the codebase that allows to easily train deep neural network based automatic speech recognition systems and run them within YARP. As an additional contribution we provide tools to experiment with recognition systems that are inspired by recent influential theories of speech perception and with systems that partly mimic the learning setting of an infant who has to learn the basic speech units of a language.

## CONFLICT OF INTEREST STATEMENT

The authors declare that the research was conducted in the absence of any commercial or financial relationships that could be construed as a potential conflict of interest.

## AUTHOR CONTRIBUTIONS

- Conceived and designed the ASR systems: LB, GM, BH, AM.
- Wrote the code: LB, BH, AM.
- Wrote the paper: BH, LB.

## FUNDING

The authors acknowledge the support of the European Commission project POETICON++ (grant agreement No. 288382) and ECOMODE (grant agreement No. 644096).

## REFERENCES

- Badino, L. (2016). Phonetic context embeddings for dnn-hmm phone recognition. In *Proc. of Interspeech* (San Francisco, CA, USA)
- Badino, L., Canevari, C., Fadiga, L., and Metta, G. (2012). Deep-level acoustic-to-articulatory mapping for dbn-hmm based phone recognition. In *Proc. of IEEE SLT* (Miami, FL, USA)
- Badino, L., Canevari, C., Fadiga, L., and Metta, G. (2014). An auto-encoder based approach to unsupervised learning of subword units. In *Proc. of IEEE ICASSP* (Florence, Italy)
- Badino, L., Canevari, C., Fadiga, L., and Metta, G. (2016). Integrating articulatory data in deep neural network-based acoustic modeling. *Computer Speech and Language* 36, 173–195
- Badino, L., Mereta, A., and Rosasco, L. (2015). Discovering discrete subword units with binarized autoencoders and hidden-markov-model encoders. In *Proc. of Interspeech* (Dresden, Germany)
- Caruana, R. (1997). Multitask learning. *Machine learning* 28, 41–75
- Dahl, G., Yu, D., Deng, L., and Acero, A. (2012). Context-dependent pre-trained deep neural networks for large-vocabulary speech recognition. *IEEE Transactions on Audio, Speech, and Language Processing* doi:10.1109/TASL.2011.2134090
- Garofolo, J., Graff, D., Paul, D., and Pallett, D. (1993a). CSR-I (WSJ0) Complete LDC93s6a. *Web Download. Philadelphia: Linguistic Data Consortium*
- Garofolo, J. S., Lamel, L. F., Fisher, W. M., Fiscus, J. G., Pallett, D. S., Dahlgren, N. L., et al. (1993b). TIMIT Acoustic-Phonetic Continuous Speech Corpus LDC93s1. *Web Download. Philadelphia: Linguistic Data Consortium*
- Hinton, G., Osindero, S., and Teh, Y. (2006). A fast learning algorithm for deep belief nets. *Neural Computation* 18, 1527–1554
- Lungarella, M., Metta, G., Pfeifer, R., and Sandini, G. (2003). Developmental robotics: a survey. *Connection Science* 15, 151–190
- Martín Abadi, Ashish Agarwal, Paul Barham, Eugene Brevdo, Zhifeng Chen, Craig Citro, et al. (2015). *TensorFlow: Large-Scale Machine Learning on Heterogeneous Systems*
- Pulvermüller, F. and Fadiga, L. (2010). Active perception: sensorimotor circuits as a cortical basis for language. *Nature Reviews Neuroscience* 11, 351–360
- Richmond, K., Hoole, P., and King, S. (2011). Announcing the electromagnetic articulography (day 1) subset of the mngu0 articulatory corpus. In *Proc. of Interspeech* (Florence, Italy)
- Seltzer, M. and Droppo, J. (2013). Multi-task learning in deep neural networks for improved phoneme recognition. In *Proc. of ICASSP* (Vancouver, British Columbia, Canada)
- Seltzer, M., Yu, D., and Wan, Y. (2013). An investigation of deep neural networks for noise robust speech recognition. In *Proc. of ICASSP* (Vancouver, British Columbia, Canada)
- Versteegh, M., Thiollie, R., Schatz, T., Cao, X. N., Anguera, X., Jansen, A., et al. (2015). The zero resource speech challenge 2015. In *Proc. of Interspeech* (Dresden, Germany)
- Vertanen, K. (2006). *Baseline WSJ Acoustic Models for HTK and Sphinx: Training Recipes and Recognition Experiments*. Tech. rep., Cavendish Laboratory

- 312 Vincent, E., Watanabe, S., Nugraha, A. A., Barker, J., and Marxer, R. (2016). An analysis of environment,  
313 microphone and data simulation mismatches in robust speech recognition. *Computer Speech & Language*  
314 doi:10.1016/j.csl.2016.11.005
- 315 Woodland, P., Odell, J., Valtchev, V., and Young, S. (1994). Large vocabulary continuous speech recognition  
316 using HTK (IEEE), vol. ii, II/125–II/128. doi:10.1109/ICASSP.1994.389562
- 317 Young, S., Evermann, G., Gales, M., Hain, T., Kershaw, D., Liu, X., et al. (2015). *The HTK book (for HTK*  
318 *version 3.5)* (Cambridge University Engineering Department)